

WEEK 6 — MACHINE LEARNING ENGINEERING

(Data Science + Feature Engineering + Model Building + Deployment + Monitoring)

WEEK 6 OBJECTIVES

Interns will learn:

Professional ML pipeline architecture

Clean, modular ML system design (Data → Features → Model → Evaluation → Deployment)

Advanced feature engineering + selection techniques

Model optimization + regularization + hyperparameter tuning

Training pipeline orchestration

Model evaluation + explainability

Model deployment + monitoring

This week produces engineers who can work on **production ML systems**.

DAY 1 — DATA PIPELINE + EDA + PROJECT ARCHITECTURE

- ◆ **Learning Outcomes:**

ML project architecture

Dataset versioning

Exploratory Data Analysis (EDA)

Professional folder structure

Data preprocessing pipelines

- ◆ **Mandatory Folder Structure (No shortcuts)**

```
src/
  data/
    raw/
    processed/
    external/
  notebooks/
  features/
  pipelines/
```

```
models/  
training/  
evaluation/  
deployment/  
monitoring/  
utils/  
config/  
logs/
```

- ◆ **Topics to Learn**

- Train/validation/test splitting strategy
- Handling missing values (mean/median/interpolate)
- Outlier detection (Z-score / IQR)
- Data scaling (StandardScaler / MinMaxScaler)
- Class imbalance (SMOTE/weights)
- Dataset versioning (DVC / folder hashing)

- ◆ **Exercise**

Build a **data loader + EDA pipeline** that:

1. Loads dataset from `/data/raw`
2. Cleans data (missing, duplicates, outliers)
3. Creates `/data/processed/final.csv`
4. Generates EDA report:
 - Correlation matrix
 - Feature distributions
 - Target distribution

- Missing values heatmap

- ✓ Data loaded
- ✓ Cleaned dataset saved
- ✓ EDA report generated

- ◆ **Deliverables**

/pipelines/data_pipeline.py
/notebooks/EDA.ipynb
/data/processed/final.csv
DATA-REPORT.md

DAY 2 — FEATURE ENGINEERING + FEATURE SELECTION

- ◆ **Learning Outcomes:**

Create meaningful features
Encoding strategies
Advanced transformations
Feature selection techniques

- ◆ **Topics**

- Categorical encoding (OneHot, Target, Label)
- Numerical feature transformations (log, sqrt, power)
- Date/time feature extraction
- Text vectorization (TF-IDF / embeddings)
- Feature selection:
 - Correlation threshold
 - Mutual information
 - Recursive Feature Elimination (RFE)

- ◆ **Exercise**

Build a feature engineering pipeline that:

- Encodes categorical features
- Normalizes numerical features
- Generates:
 - 10+ new features
- Applies feature selection
- Produces `X_train`, `X_test`, `y_train`, `y_test`

- ✓ Feature pipeline saved
- ✓ Feature importance plotted

- ◆ **Deliverables**

```
/features/build_features.py  
/features/feature_selector.py  
/features/feature_list.json  
FEATURE-ENGINEERING-DOC.md
```

DAY 3 — MODEL BUILDING + ADVANCED TRAINING PIPELINE

- ◆ **Learning Outcomes:**

Multi-model training
Pipeline building
Cross-validation
Overfitting control

- ◆ **Topics**

- Models:

- Logistic Regression
- Random Forest
- XGBoost / LightGBM
- Neural Network
- Cross-validation (k-fold)
- Overfitting vs underfitting
- Regularization (L1/L2)
- Model comparison

◆ **Exercise**

Create a unified training pipeline that:

- Trains 4 models
- Performs 5-fold cross-validation
- Outputs:
 - Accuracy
 - Precision
 - Recall
 - F1 Score
 - ROC-AUC

Save the best model automatically.

- ✓ Best model selected
- ✓ Metrics saved
- ✓ Confusion matrix plotted

◆ **Deliverables**

/training/train.py
/models/best_model.pkl
/evaluation/metrics.json
MODEL-COMPARISON.md

DAY 4 — HYPERPARAMETER TUNING + EXPLAINABILITY + ERROR ANALYSIS

◆ Learning Outcomes:

Optimize a model like an ML engineer
Interpret model decisions
Perform deep error analysis

◆ Topics

- GridSearch / RandomSearch / Bayesian tuning (Optuna)
- SHAP values
- Feature importance
- Error clustering
- Bias/variance analysis

◆ Exercise

Implement:

- Hyperparameter tuning with Optuna/GridSearch
- Generate:
 - SHAP summary plot
 - Feature importance chart
 - Error analysis heatmap

- ✓ Improvement over baseline
- ✓ Explainability added

- ◆ **Deliverables**

/training/tuning.py
/evaluation/shap_analysis.py
/tuning/results.json
MODEL-INTERPRETATION.md

DAY 5 — MODEL DEPLOYMENT + MONITORING + MLOPS CONCEPTS (CAPSTONE)

- ◆ **Learning Outcomes:**

Deploy real ML systems
Monitor drift
Production-ready pipelines

- ◆ **Topics**

- Serving model using FastAPI/Flask
- Creating `/predict` endpoint
- Saving & loading with joblib/pickle
- Input schema validation
- Monitoring:
 - Data drift
 - Accuracy decay
- Logging predictions

- ◆ **Exercise (Capstone)**

Deploy model as an API:

```
POST /predict
{
  "feature1": value,
  "feature2": value,
  ...
}
```

Add:

- ✓ Prediction logging
- ✓ Request ID tracking
- ✓ Input validation
- ✓ Versioned model loading
- ✓ Basic dashboard (Streamlit optional)

Prepare:

- Dockerfile
- requirements.txt
- .env.example
- README.md

◆ **Deliverables**

```
/deployment/api.py
/deployment/Dockerfile
/monitoring/drift_checker.py
/prediction_logs.csv
DEPLOYMENT-NOTES.md
```

WEEK- 6 COMPLETION REQUIREMENTS

Skill Area	Requirement
Data Engineering	Clean, versioned datasets

Feature Engineering	Automated + selected features
Model Building	Multiple models trained
Optimization	Hyperparameter tuning
Explainability	SHAP + feature importance
Deployment	Working API
Monitoring	Drift detection
Documentation	Full reports

EXPECTED OUTCOME AFTER WEEK- 6

Interns can now:

- ✓ Design end-to-end ML pipelines
 - ✓ Engineer advanced features
 - ✓ Train, tune and compare models
 - ✓ Explain model behaviour
 - ✓ Deploy ML systems as APIs
 - ✓ Monitor model performance
 - ✓ Work in real-world ML teams
-