

# WEEK 7 — GENAI & MULTIMODAL RAG ENGINEERING

(*Text RAG + Image RAG + SQL Question Answering + Hybrid Retrieval + Local LLMs*)

## **USE CASE: Enterprise Knowledge Intelligence System (Text + Images + SQL RAG)**

Interns will build an enterprise-grade GenAI system that can:

- Answer questions from internal documents (PDFs, DOCX, TXT)
- Retrieve & reason over images (diagrams, charts, forms)
- Query structured data using natural language → SQL → result
- Run on **either** local open-source models **or** hosted LLM APIs (OpenAI / Claude / Gemini)
- Provide faithful, non-hallucinated responses with evaluation & monitoring

This simulates real enterprise systems used in:

- Banking (policy manuals + scanned KYC files + SQL transaction DB)
- Insurance (claims PDFs + damage images + SQL claim tables)
- Manufacturing (blueprints + manuals + operational DB)

---

## **DATASETS FOR THIS USE-CASE (WORKING + RELEVANT)**

### **1) Tabular / Structured Data (CSV – Enterprise)**

**Enterprise CSV Samples (General Business Tables)**

- Type: Tabular CSV files for business data (employees, products, transactions, etc.)
- Source: <https://www.datablist.com/learn/csv/download-sample-csv-files>

## 2) Document / Knowledge Data (Text-RAG)

### Enterprise RAG Markdown Dataset

- Type: Markdown text documents — suitable for ingestion & retrieval tasks
- Source: <https://www.kaggle.com/datasets/rrr3try/enterprise-rag-markdown>

## 3) Graph / Structured Relationship Data

### Graphs Dataset

- Type: Structured graph data (networks/relationships) — useful for entity graphs, KG-RAG
  - Source: <https://www.kaggle.com/datasets/sunedition/graphs-dataset>
- 

## WEEK 7 OBJECTIVES

Interns will learn:

- ✓ Advanced Retrieval-Augmented Generation (RAG) architecture
  - ✓ Local LLM setup & optimization **or** API-based LLM integration
  - ✓ Hybrid retrieval (semantic + keyword + reranking + image)
  - ✓ Multimodal embeddings for Image-RAG
  - ✓ SQL-based question answering
  - ✓ Prompt + context optimization
  - ✓ Document chunking & metadata routing
  - ✓ GenAI pipeline design (fully offline **or** “enterprise-controlled API mode”)
- 

## MODEL STACK (CHOOSE ONE PATH — OPEN-SOURCE LOCAL OR HOSTED APIS)

## **PATH A — OPEN-SOURCE LOCAL (Fully Offline)**

### **LLMs (choose 1):**

- Mistral-7B Instruct
- LLaMA-2 / LLaMA-3 local versions
- Qwen2
- Phi-3

### **Embeddings Models:**

- BGE-small / BGE-base
- Instructor-XL
- GTE-base
- CLIP (for image embeddings)

### **Vector DB Options:**

- FAISS
- Chroma
- Qdrant (local container)

### **OCR / Image Tools:**

- Tesseract
- OpenCV
- BLIP / CLIP

### **SQL Database:**

- PostgreSQL

- SQLite (recommended for interns)
- 

## PATH B — HOSTED LLMs (Keys + Model IDs)

Use this when interns don't have GPUs or you want "enterprise-grade APIs" while keeping the *same* RAG architecture, chunking, tracing, eval, and monitoring.

### Option 1: OpenAI (API)

**API Key env var:** `OPENAI_API_KEY`

**Common model choices:** GPT-5.2 / GPT-5 mini / GPT-4.1 (and related variants). [OpenAI Platform+2OpenAI+2](#)

### Option 2: Anthropic Claude (API)

**API Key env var:** `ANTHROPIC_API_KEY`

**Common model choices:** Claude 3.7 Sonnet (and other Claude family models listed in Claude docs). [Anthropic+1](#)

### Option 3: Google Gemini (API)

**API Key env var:** `GOOGLE_API_KEY` (*Gemini API / AI Studio*)

**Common model choices:** Gemini 3 Pro / Gemini Flash family (per Gemini models docs). [Google AI for Developers+1](#)

**Important:** Path B is *not* fully offline. For enterprise realism: add policy gates (PII redaction), logging controls, and strict "no training on customer data" vendor settings per your org's policy.

---

## Minimal “Provider Switch” Config (keep your pipeline identical)

Create `config/model.yaml`:

- `provider: local | openai | anthropic | gemini`
- `model_name: <chosen model id>`
- `api_key_env: OPENAI_API_KEY | ANTHROPIC_API_KEY | GOOGLE_API_KEY`

Your code only changes in **one place**: `/generator/llm_client.py` (local loader vs API client). Everything else stays the same.

---

# DAY 1 — LOCAL RAG SYSTEM + PIPELINE ARCHITECTURE

## Learning Outcomes

- RAG architecture (Retriever → Generator)
- Local LLM loading and inference
- Embedding generation
- Document chunking strategies
- Semantic indexing (HNSW / IVF / Flat)

## Mandatory Folder Structure

```
src/
  data/
    raw/
    cleaned/
    chunks/
  embeddings/
  vectorstore/
  retriever/
  generator/
  pipelines/
  prompts/
  models/
  evaluation/
  utils/
  config/
  logs/
```

## Topics to Learn

- RAG architecture fundamentals
- Chunk size vs. token limits
- Overlap strategy
- Embedding pipelines
- Metadata tagging
- Vector index structures

## Exercise

Build a local ingestion & chunking pipeline:

Pipeline must:

- Load PDFs / TXT / CSV / DOCX
- Clean text → split into 500–800 token chunks
- Add metadata (source, page number, tags)
- Generate local embeddings
- Store vectors in FAISS or Qdrant
- Build retriever module

- ✓ Documents loaded
- ✓ Chunks created
- ✓ Embeddings generated
- ✓ Vector DB initialized

## Deliverables

- `/pipelines/ingest.py`
- `/embeddings/embedder.py`

- `/vectorstore/index.faiss`
  - `/retriever/query_engine.py`
  - `RAG-ARCHITECTURE.md`
- 

## DAY 2 — ADVANCED RETRIEVAL + CONTEXT ENGINEERING

### Learning Outcomes

- Improve retrieval accuracy
- Build hybrid retrieval (semantic + keyword + reranking)
- Context ranking strategies
- Reduce hallucination

### Topics

- Hybrid search: BM25 + embeddings
- Reranking (cross-encoder / cosine)
- Max Marginal Relevance (MMR)
- Chunk deduplication
- Optimizing context window for LLMs

### Exercise

Build an advanced retriever that supports:

```
query = "Explain how credit underwriting works"
```

```
top_k = 5
filters = { "year": "2024", "type": "policy" }
```

Add:

- Keyword fallback
- Reranking
- Deduplication
- Traceable context sources

- ✓ Higher precision
- ✓ Lower hallucination
- ✓ Fully traceable context

## Deliverables

- `/retriever/hybrid_retriever.py`
  - `/retriever/reranker.py`
  - `/pipelines/context_builder.py`
  - `RETRIEVAL-STRATEGIES.md`
- 

# DAY 3 — IMAGE-RAG (MULTIMODAL RAG)

## Learning Outcomes

- Handle images inside RAG
- Generate & store vision embeddings
- OCR extraction + captioning

- Image similarity search

## Topics

- CLIP embeddings (image + text)
- OCR extraction using Tesseract
- Caption generation using BLIP
- Multimodal vector DB design

## Exercise

Build an Image RAG pipeline:

Supports ingestion of:

- PNG, JPG, scanned PDFs, forms, diagrams

System generates:

- OCR Text
- CLIP embeddings
- Captions (BLIP)
- Multimodal vector index

Query modes:

- Text → Image
- Image → Image
- Image → Text Answer

Example:

User uploads an engineering diagram → System retrieves related diagrams + explanations.

## **Deliverables**

- `/pipelines/image_ingest.py`
  - `/embeddings/clip_embedder.py`
  - `/retriever/image_search.py`
  - `MULTIMODAL-RAG.md`
- 

# **DAY 4 — SQL QUESTION ANSWERING SYSTEM (Text → SQL → Answer)**

## **Learning Outcomes**

- Convert natural language queries into SQL
- Schema-aware reasoning
- SQL query correction / validation
- Injection-safe execution

## **Topics**

- Schema extraction
- Prompting patterns for SQL generation
- Query validation
- Error correction
- Summarizing result tables

## Exercise

Build a SQL-QA Engine:

User: "Show total sales by artist for 2023."

System should:

- Generate SQL using LLM
- Validate SQL
- Execute on SQLite/PostgreSQL
- Summarize the results

Features:

- ✓ Auto schema loader
- ✓ Query validator
- ✓ Safe executor
- ✓ Result summarizer

## Deliverables

- `/pipelines/sql_pipeline.py`
  - `/generator/sql_generator.py`
  - `/utils/schema_loader.py`
  - `SQL-QA-DOC.md`
- 

# DAY 5 — ADVANCED RAG + MEMORY + EVALUATION (CAPSTONE)

## Learning Outcomes

- Add conversational memory
- Self-reflection + refinement loops
- Hallucination detection
- Faithfulness scoring
- Production-ready API structure

## Topics

- Memory storage (Vector + Redis + Local File)
- Self-critique for improving answers
- Context match score
- Faithfulness scoring
- Human feedback logging

## Exercise (Capstone)

Build a complete system with endpoints:

- `/ask`
- `/ask-image`
- `/ask-sql`

Add:

- ✓ Memory for last 5 messages
- ✓ Refinement loop
- ✓ Hallucination detection
- ✓ Confidence score
- ✓ Logging + debugging traces
- ✓ Streamlit UI or CLI

## Deliverables

- `/deployment/app.py`
  - `/evaluation/rag_eval.py`
  - `/memory/memory_store.py`
  - `CHAT-LOGS.json`
  - `DEPLOYMENT-NOTES.md`
- 

## WEEK 7 COMPLETION REQUIREMENTS

Skill Area	Requirement
RAG System	End-to-end functioning
Advanced Retrieval	Hybrid + reranking
Image RAG	CLIP + OCR + multimodal vectors
SQL QA	Natural language → SQL → Answer
Memory	Context retention
Evaluation	Faithfulness & context matching
Local Models / API Models	Works with chosen stack
Documentation	All systems documented

---

## EXPECTED OUTCOME

After this week, interns will be able to:

- ✓ Build enterprise-grade GenAI systems
- ✓ Handle documents + images + SQL databases
- ✓ Run LLMs locally **or** via OpenAI/Claude/Gemini keys
- ✓ Reduce hallucination by over 70%
- ✓ Build multimodal & hybrid RAG pipelines
- ✓ Prepare for multi-agent GenAI scaling

